
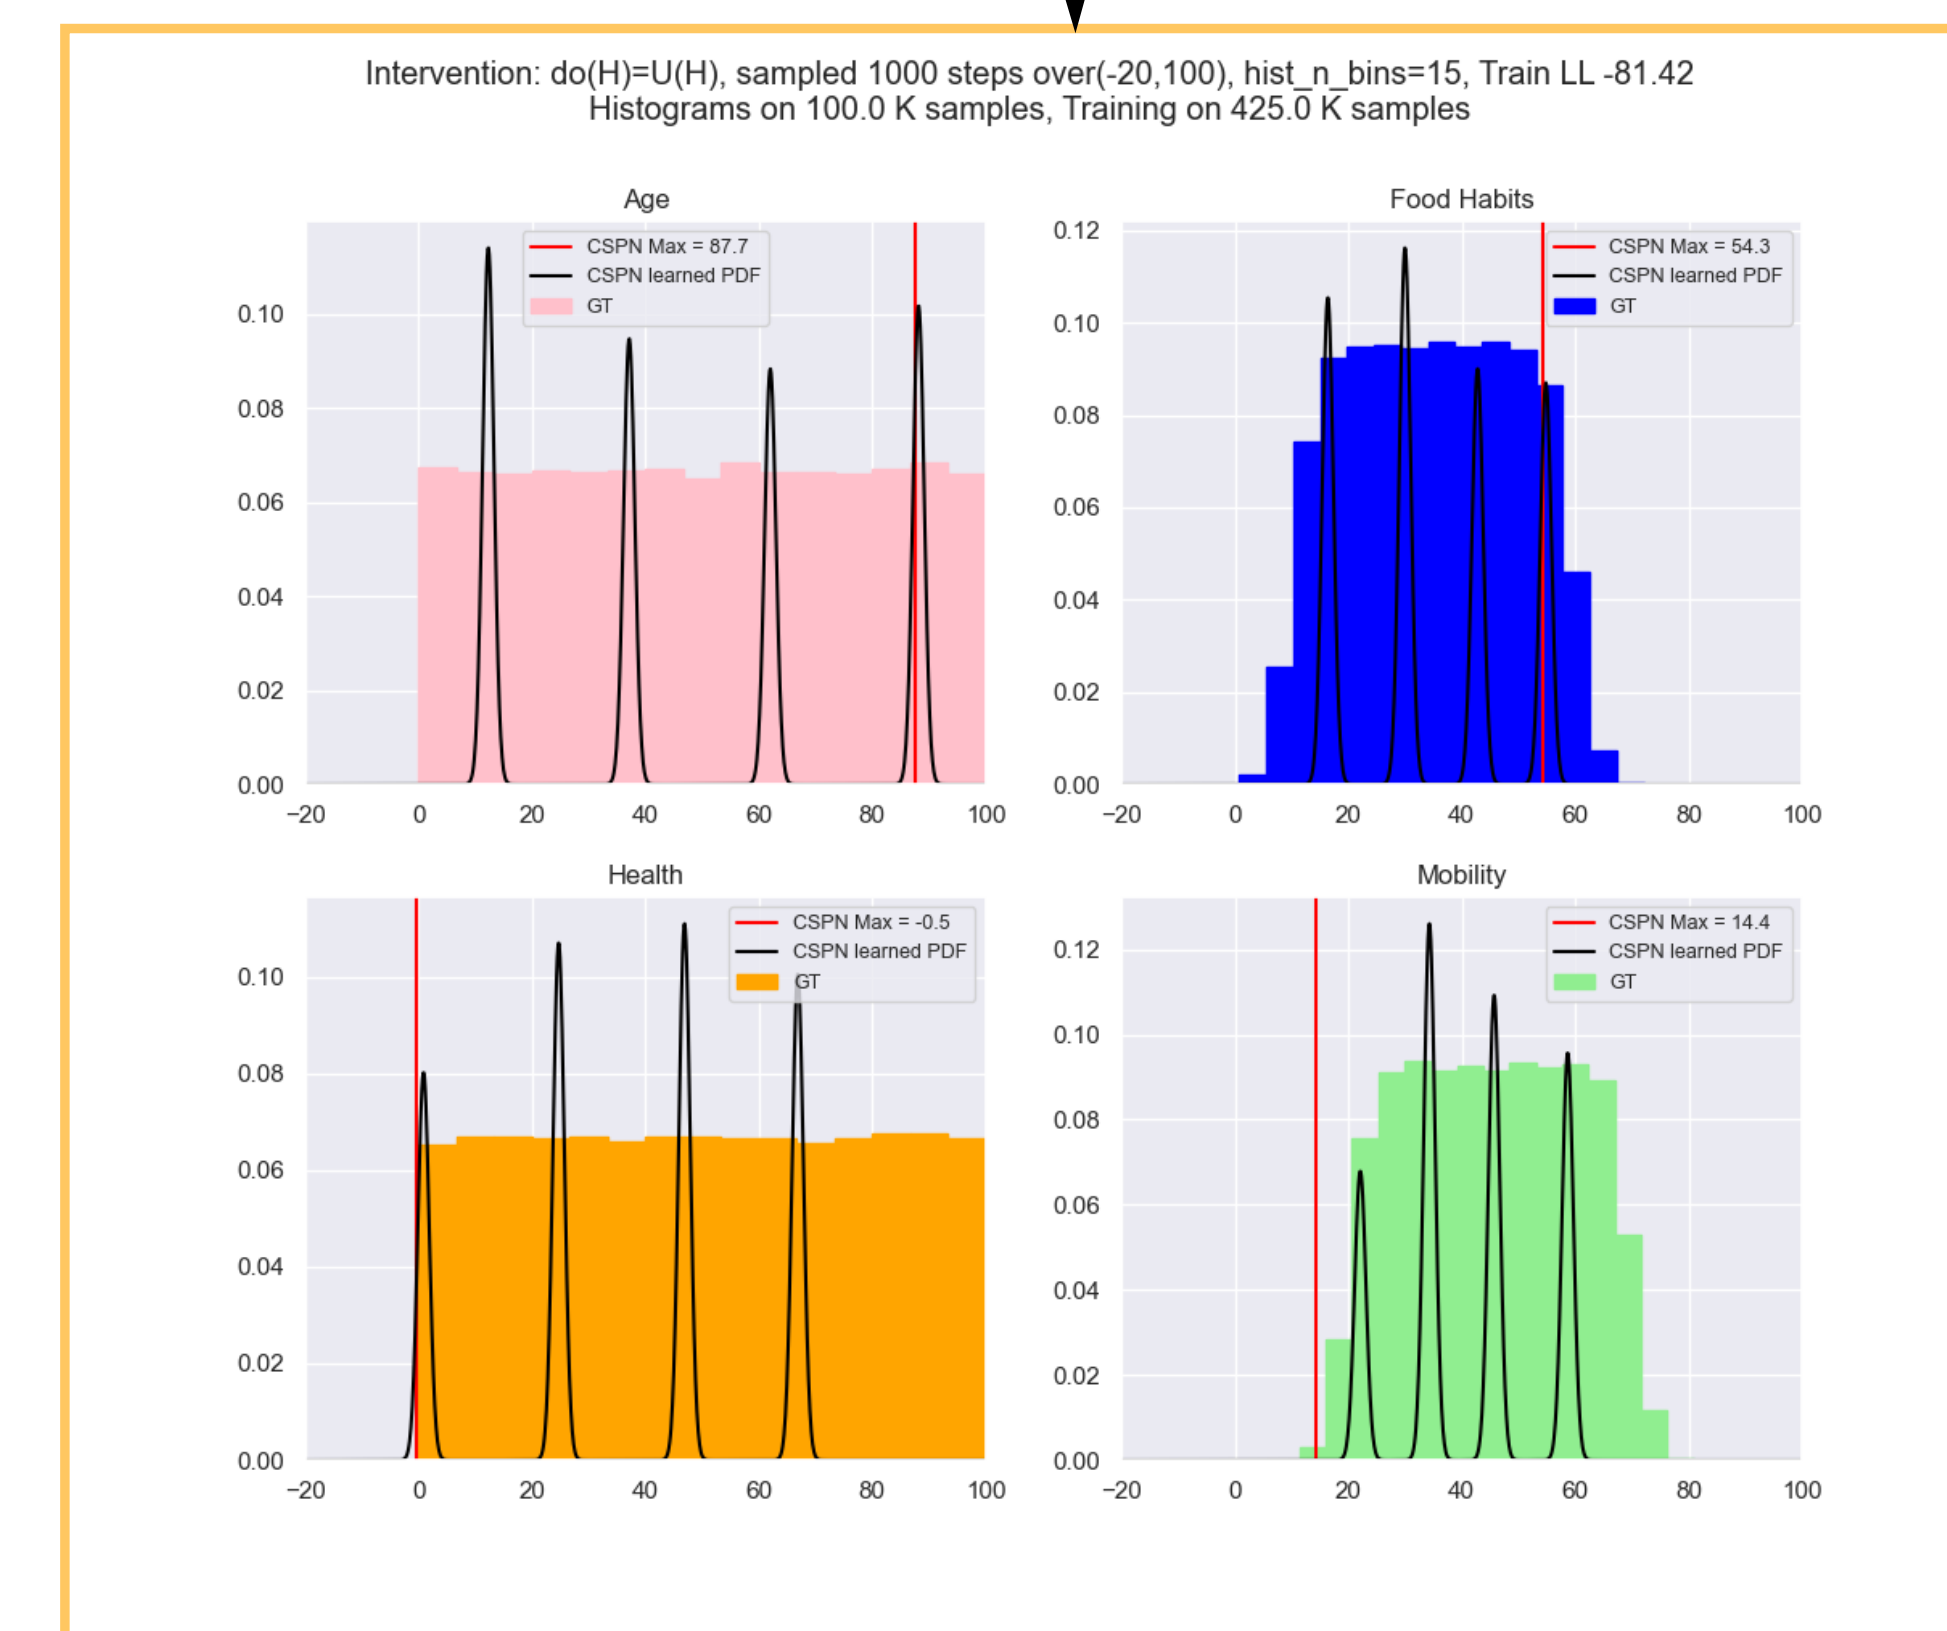

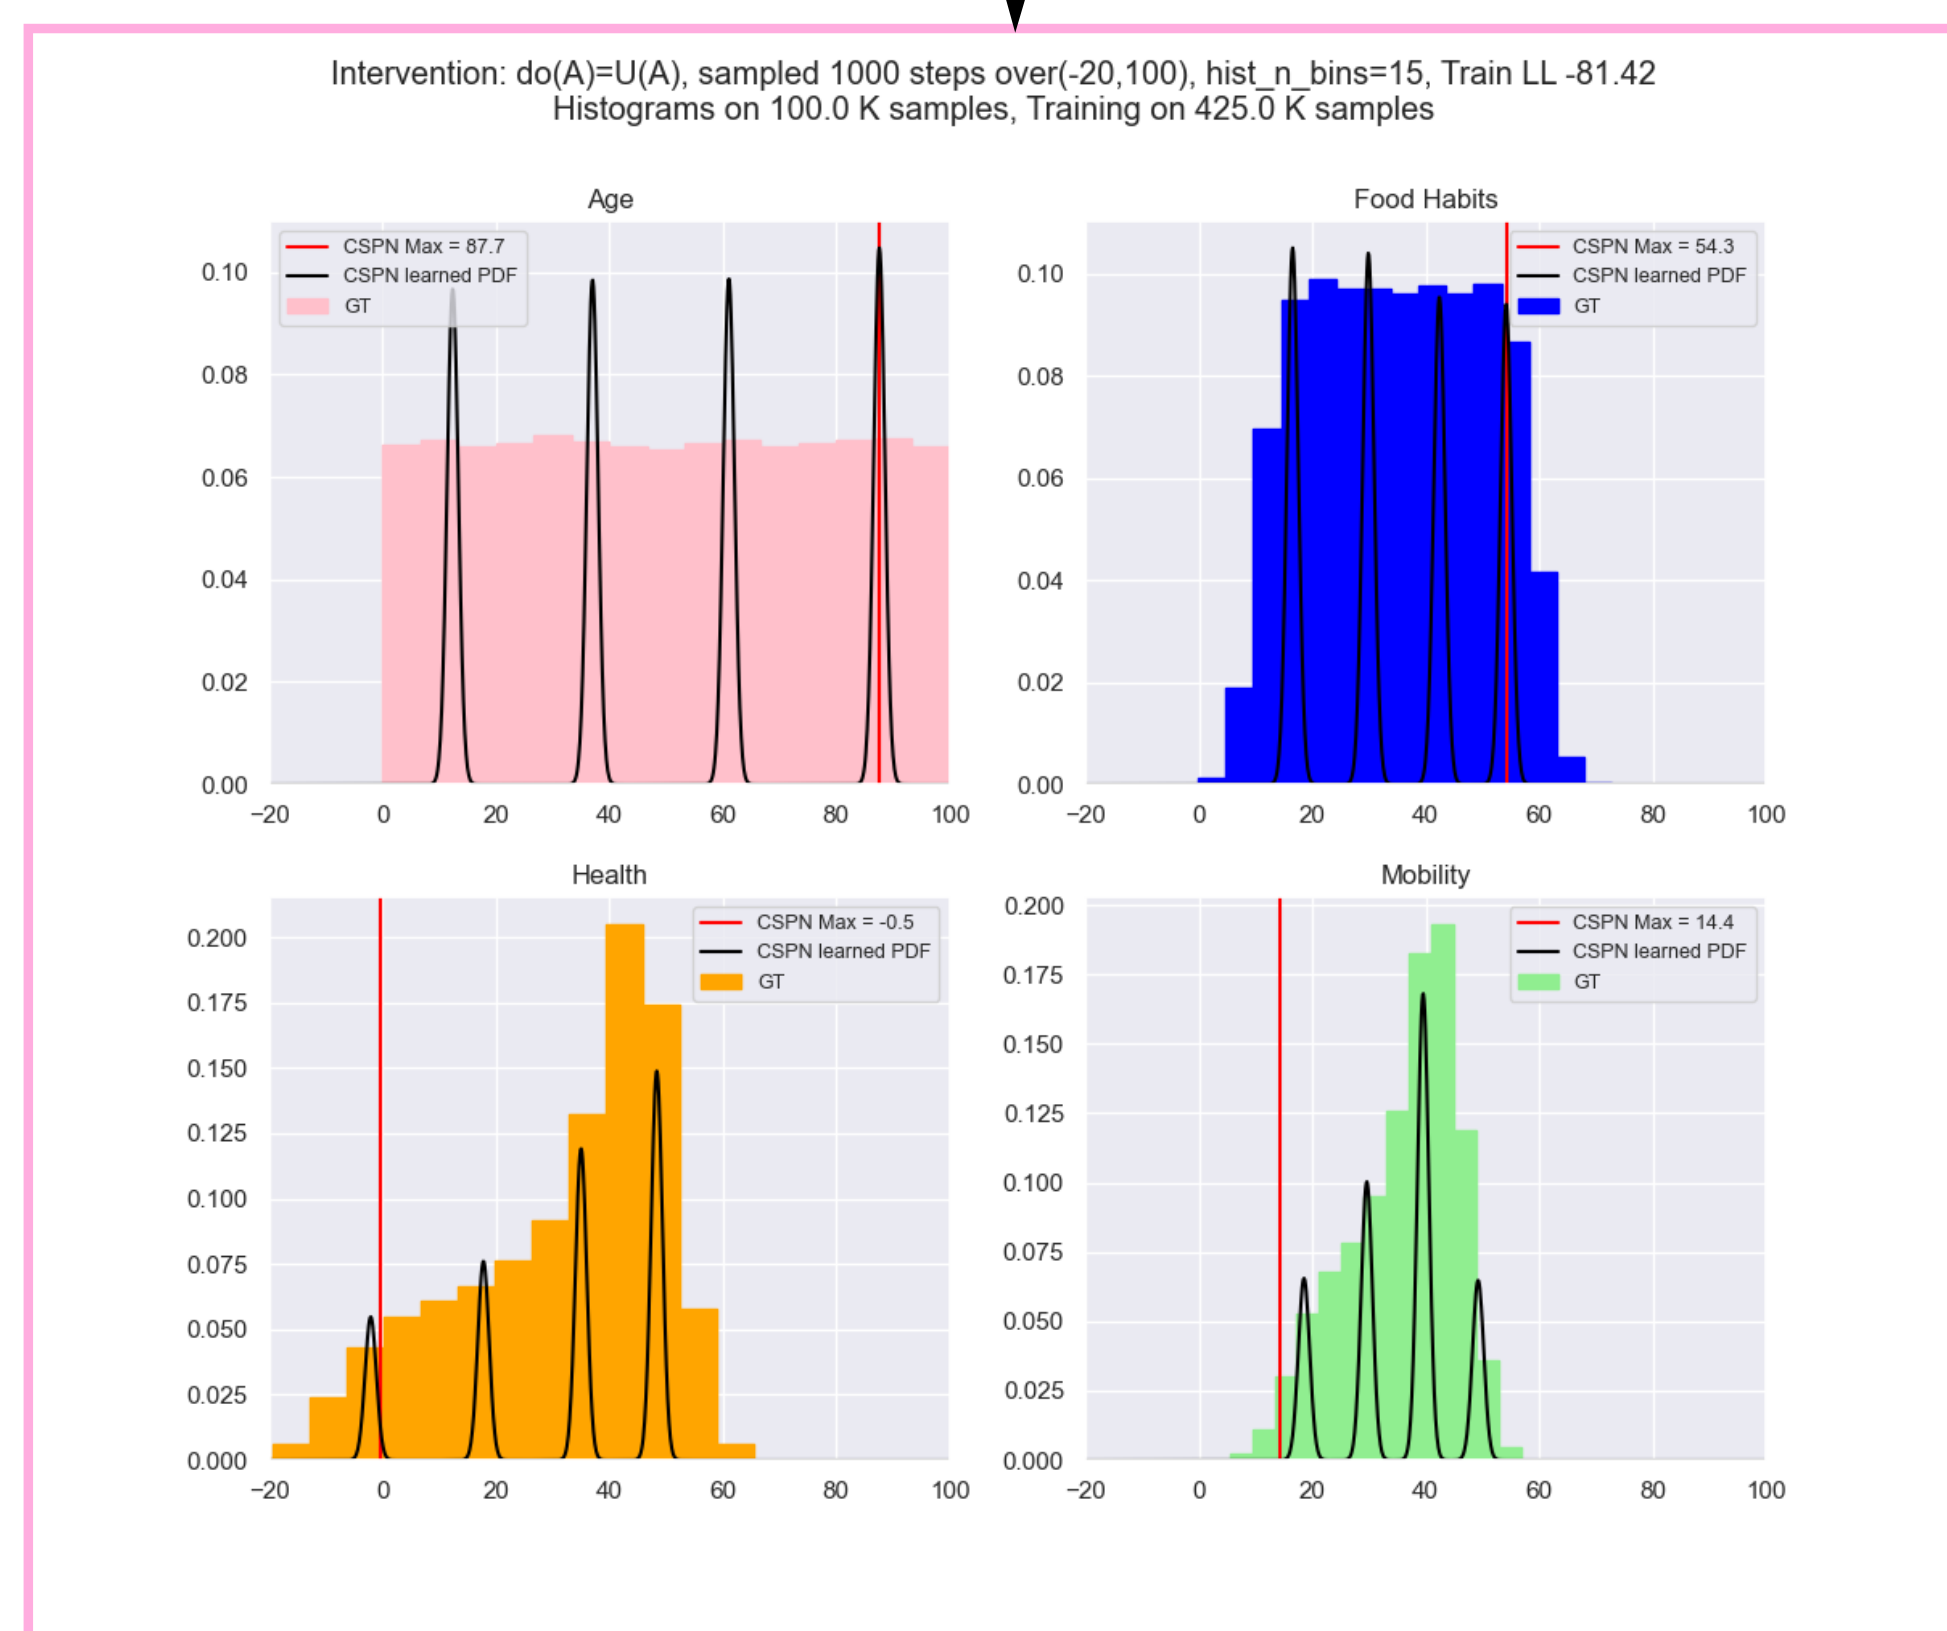
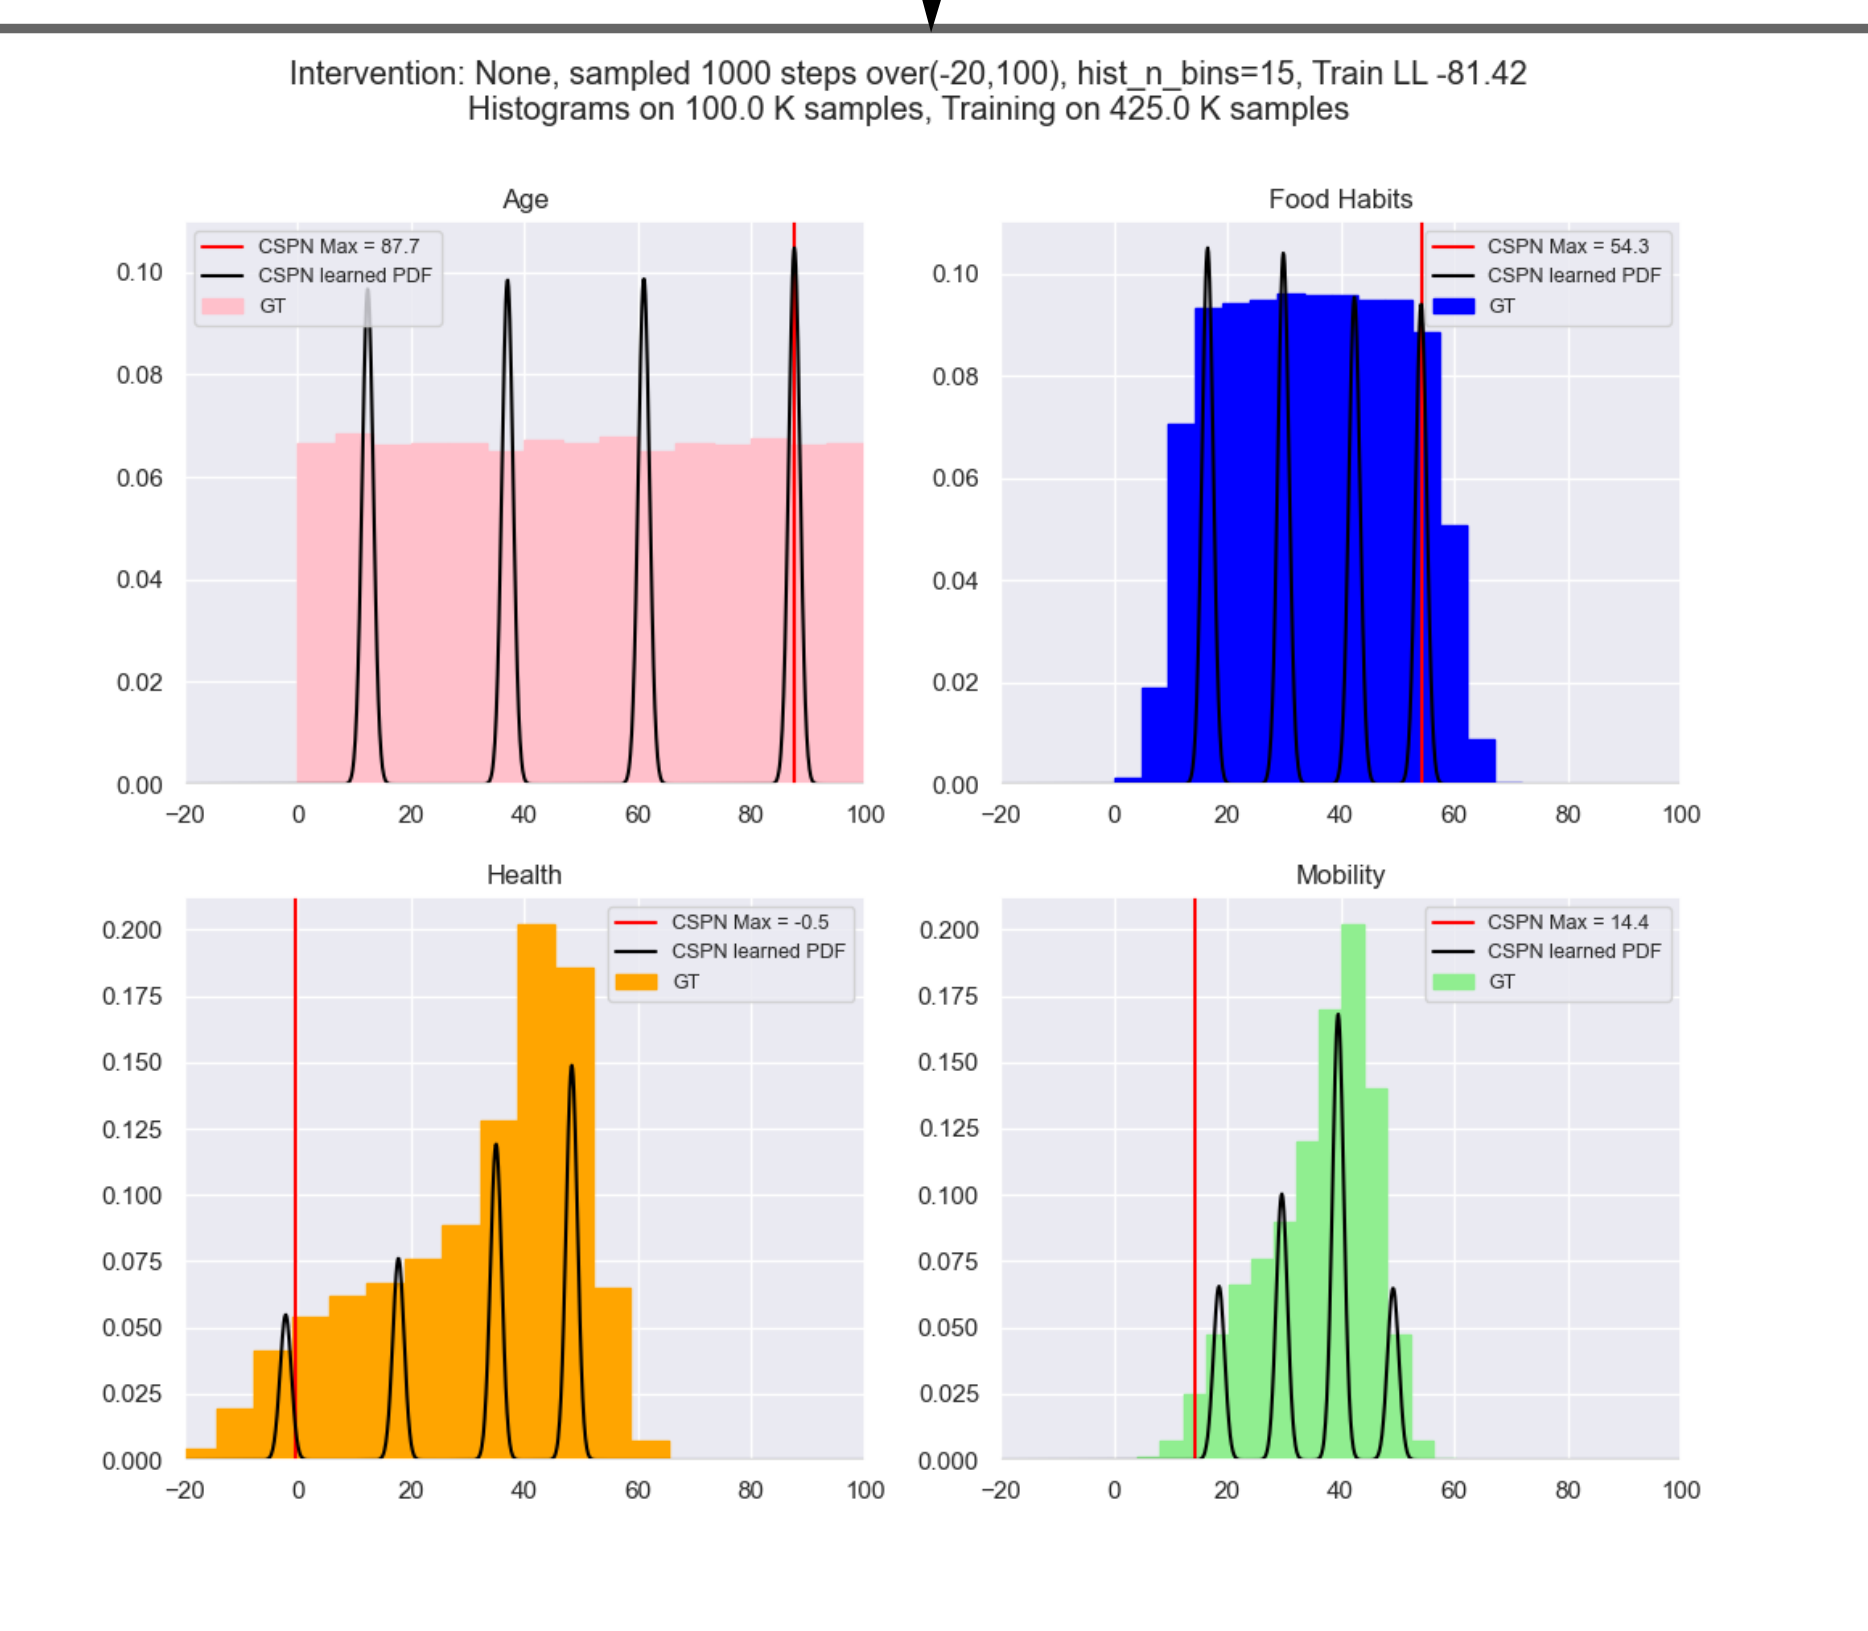
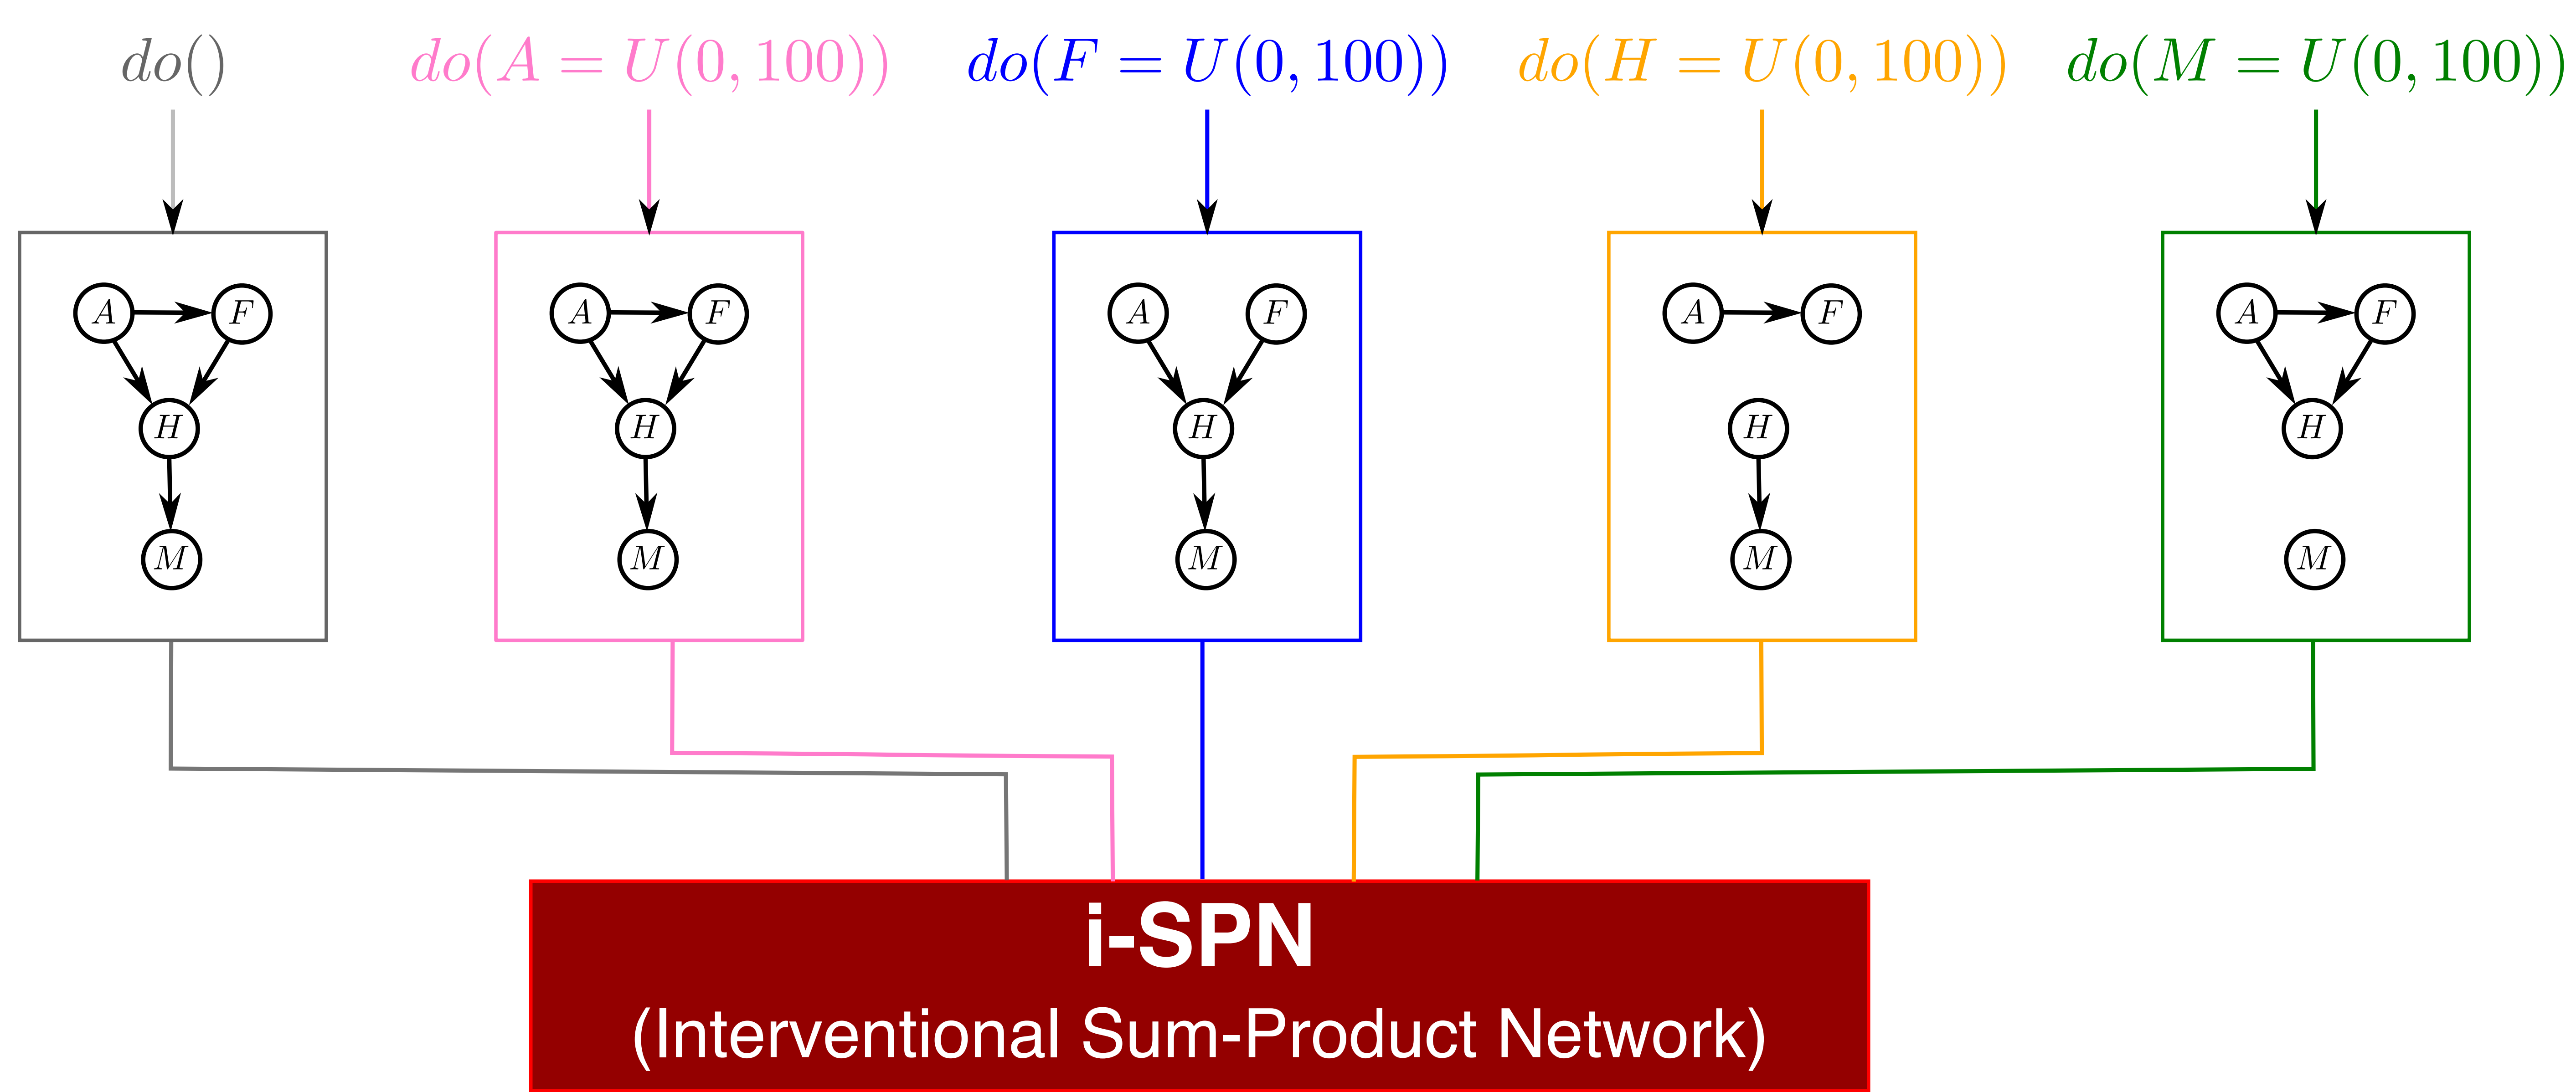
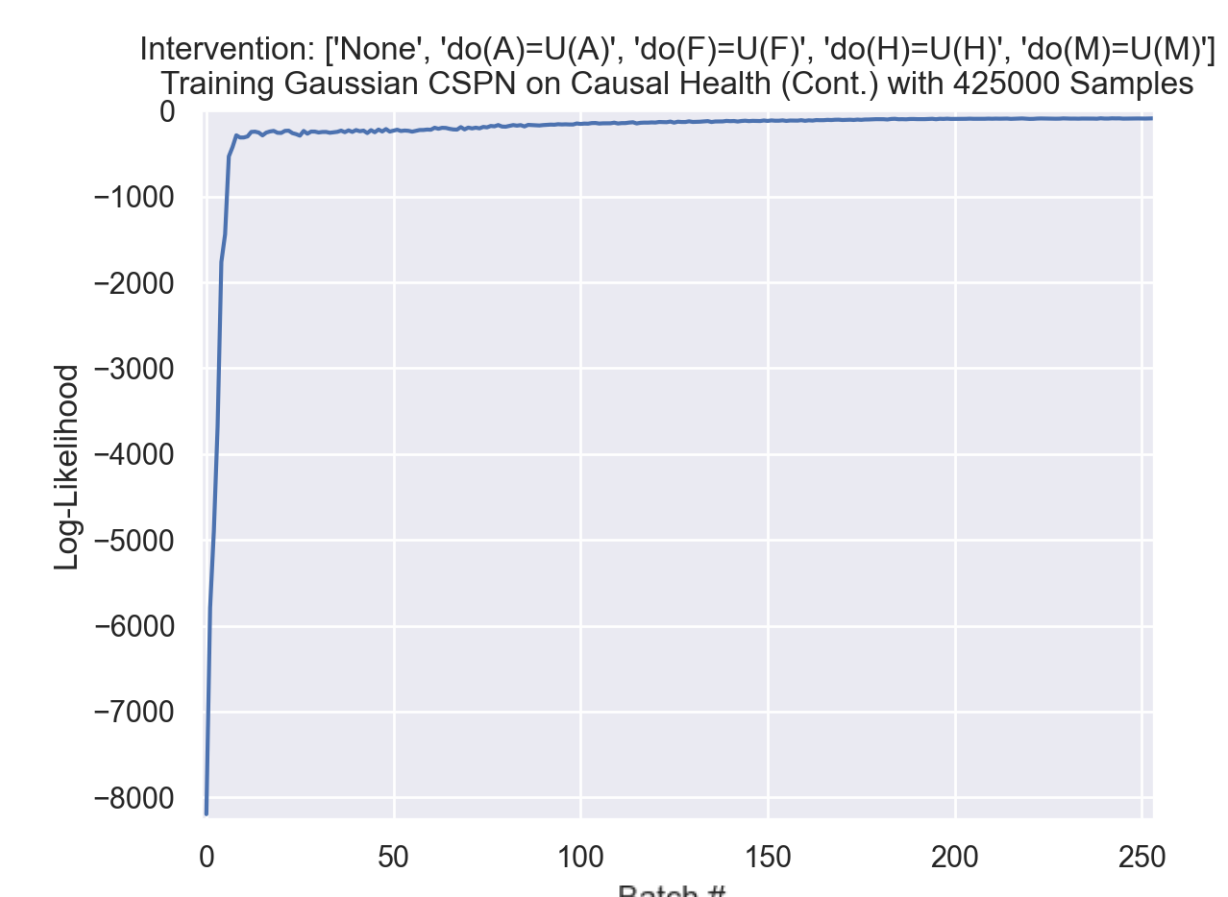


Results for the Causal Health Dataset
4 variables, mixed distributions, continuous
Ground Truth Histograms vs. Learned PDF
 Input: flattened Graph Adjacency Matrices
 Training Log-Likelihood -81.4
 Training on 425,000 samples for 130 Epochs
 Batch Size 1,000, # Sums 4, # Gaussians 4
 Total # Parameters 16,304
 Reproducible across multiple seeds



Results for the ASIA Dataset
8 variables, Bernoulli, discrete
Ground Truth Histograms vs. Learned PDF
 Input: flattened Graph Adjacency Matrices
 Training Log-Likelihood -0.87
 Training on 42,500 samples for 20 Epochs
 Batch Size 100, # Sums 4, # Gaussians 4
 Total # Parameters 11,080
 Reproducible across multiple seeds

